

Multiway Classification with Real Data

Matthew Asisgress

2025-10-29

In this demonstration, we apply R package **cpfa** (Asisgress, 2025) to three real datasets. First, we apply the package to the MNIST dataset (LeCun, Cortes, and Burges, 1998; LeCun et al., 2002)—showing how to use the package to distinguish between digits of 2 and 3, a binary classification problem. We further apply the package to the Fashion MNIST dataset (Xiao, Rasul, and Vollgraf, 2017)—distinguishing among images of tops, trousers, and sandals, which is a multiclass classification problem. Finally, we apply the package to the VesselMNIST3D dataset (Yang et al., 2020) from the MedMNIST database (Yang et al., 2023; Yang, Shi, Ni, 2021), which contains three-dimensional representations of blood vessels. Specifically, we distinguish healthy blood vessels from aneurysm blood vessels, which is a binary classification problem.

A.) MNIST

The MNIST dataset consists of 70,000 grayscale images of handwritten digits from 0 to 9. We use R package **dslabs** (Irizarry and Gill, 2025) to download the dataset. We subset to only 2000 images from the original training set. The images include 980 images of the digit 2 and 1020 images of the digit 3.

Download

```
# load library
library(dslabs)

# download MNIST data and subset training set to digits of 2 or 3
mnist <- read_mnist()
inde <- which(mnist$train$labels %in% c(2, 3))
images <- mnist$train$images[inde, ]
labels <- mnist$train$labels[inde]


# restructure data into a three-way array and prepare labels
X0 <- array(images, c(nrow(images), 28, 28))
y0 <- as.factor(as.numeric(as.factor(labels)) - 1)

# subset to first 2000 images
ind <- 1:2e3
X <- X0[ind, , ]
y <- y0[ind]
```

We plot and examine an example of the digit 2 and the digit 3.

```
# plot example of digits 2 and 3
par(mfrow = c(1, 2))
pdigit <- function(imat) {
  m <- t(apply(imat, 2, rev))
  image(m, col = gray(seq(0, 1, 0.05)), xaxt = "n", yaxt = "n")
}
```

```
}
for (i in 0:1) {pdigit(t(X[which(y == i)[1], , ]))}
```



Analysis

We load R package **cpfa** and initialize the tensor model. First, the data array is a regular three-way array; so we set `model <- "parafac"` to use a Parafac model (Harshman, 1970). Second, we initialize the number of components to fit for the model by setting `nfac <- c(2, 3)` in order to fit both a two-component Parafac model and a three-component Parafac model. We set `nstart <- 10` to allow for 10 random starts in the Parafac alternating least squares algorithm fit by R package **multiway** (Helwig, 2025), upon which R package **cpfa** depends. Third, we specify the constraint desired for each array mode using `const`. Fourth, we use `cmode <- 1` to specify that the classification mode is the first mode of the input array (i.e., the mode connected to the class labels). Note that numerous constraint options are available; after loading **cpfa**, type `const()` in the R console to access a constraint options list provided through R package **CMLS** (Helwig, 2025).

Next, we initialize classification methods. First, we use `method = c("PLR", "RF")` to employ penalized logistic regression (PLR) and random forest (RF) classifiers for this problem. See `help(cpfa)` for information on additional classification methods available. Second, we specify that the problem is a binary classification problem by setting `family <- "binomial"`. Third, we use 10-fold cross-validation (CV) in our inner training, setting `nfolds <- 10`; and fourth, we set `nrep <- 5` to perform five outer train-test splits of our data. Fifth, we set `ratio <- 0.9` to specify that each outer training set contains a proportion of 0.9 of the full input data while the outer testing set contains a proportion of 0.1. Finally, we specify ranges for tuning parameters alpha (PLR), the number of trees (RF), and node size (RF), wrapping them into a list called `parameters`.

```
# load library
library(cpfa)

# set seed
set.seed(500)

# initialize model
model <- "parafac"
nfac <- c(2, 3)
nstart <- 10
const <- c("uncons", "uncons", "uncons")
cmode <- 1
```

```


# initialize classification
method <- c("PLR", "RF")
family <- "binomial"
nfolds <- 10
nrep <- 5
ratio <- 0.9

# initialize tuning parameters
alpha <- seq(0, 1, length = 8)
ntree <- c(400, 600, 800, 1000)
nodesize <- c(4, 8, 16, 32)
parameters <- list(alpha = alpha, ntree = ntree, nodesize = nodesize)

# implement train-test splits with inner k-fold CV to optimize classification
outputR <- cpfa(x = X, y = y, model = model, nfac = nfac, nstart = nstart,
               const = const, cmode = cmode, method = method, family = family,
               nfolds = nfolds, nrep = nrep, ratio = ratio,
               parameters = parameters, type.out = "descriptives",
               seeds = NULL, plot.out = TRUE, parallel = FALSE,
               verbose = FALSE)

```

Performance Measure



Results

We examine classification performance metrics of error (`err`) and overall accuracy (`acc`) for each model and for each classifier, looking at their median across outer train-test splits. We also examine, averaged across train-test splits, the optimal tuning parameters chosen (i.e., those that minimized misclassification error in the inner 10-fold CV) for each classifier.

```
# examine classification performance measures - median across train-test splits
outputR$descriptive$median[, 1:2]
```

```
##           err    acc
## fac.2plr 0.110 0.890
## fac.2rf  0.095 0.905
## fac.3plr 0.110 0.890
## fac.3rf  0.105 0.895
```

```
# examine optimal tuning parameters averaged across train-test splits
outputR$mean.opt.tune
```

```
##   nfac      alpha  lambda gamma cost ntree nodesize size decay rda.alpha delta
## 1    2 0.11428571 2.350755    NA   NA   600    32.0   NA    NA        NA    NA
## 2    3 0.02857143 6.256281    NA   NA   680    10.4   NA    NA        NA    NA
##   eta max.depth subsample nrounds
## 1  NA         NA         NA      NA
## 2  NA         NA         NA      NA
```


We can see that classification accuracy values are close to 0.9 for both two-component and three-component models. The two-component Parafac model with the RF classifier performed best. In addition, we can see average optimal tuning parameters for each classifier. Note that PLR optimized lambda internally.

Next, we use the function `plotcpfa`. The function performs two tasks: (1) it fits the optimal model among those fit using function `cpfa`; and (2) it plots the estimated component weights from that model for the non-classification modes. Looking at the plots, we can search for relationships between the levels of mode A or B and the model components. In this case, mode A corresponds to the horizontal axis of the images (e.g., such as those images shown above) while mode B corresponds to the vertical axis.


```
# set seed
set.seed(500)
```

```
# plot heatmaps of component weights for optimal model
results <- plotcpfa(outputR, nstart = 10, ctol = 1e-6, verbose = FALSE)
```

A Weights



Components B Weights



Components

For the A mode, component weights are stronger between levels 5 and 24 for the first component. Looking at the B mode, weights are stronger in the middle between levels 4 and 25. Whether viewed from the A mode or B mode perspective, this first component appears to be a general component showing the presence of non-zero values, distinguishing empty outer parts of images from non-zero inner parts. The digits 2 and 3 have systematic differences in the presence of non-zero values in these areas.

Examining the second component for the A mode, weights are stronger near and around level 14 where there appears to be a greater presence of zero values in digit 3 compared to digit 2. A similar pattern exists in the B mode around levels 18 and 19 with more non-zero values for 2, compared to 3. Taken together between the two modes, this component could be identifying the closed circle present in the digit 2, which differs from the open or half circle seen in digit 3.

B.) Fashion MNIST

The Fashion MNIST dataset consists of 70,000 grayscale images of fashion objects within 10 different categories, indexed by class labels of 0 through 9 (Xiao, Rasul, and Vollgraf, 2017). We use R package **keras3** (Kalinowski, Allaire, and Chollet, 2025) to download the dataset. We subset to only 1000 images from the training set. We only use images within the categories of top (with a label of 0), trouser (1), or sandal (5). The images include 308 tops, 357 trousers, and 335 sandals.

We remove a different number of horizontal levels from each image, randomly selecting a number of levels between three and six (inclusive) to remove for each image. The resulting array is ragged: mode A (the horizontal mode) contains a different number of rows for each image while mode B (the vertical mode) contains exactly 28 levels for all images. This removal is meant to simulate data corruption or sensor malfunction that could lead to incomplete images. To continue forward, we use a Parafac2 model (Harshman, 1972), which can be fit to a ragged array (i.e., one mode has a different number of levels, conditional on another mode). The Parafac2 model maintains properties similar to the Parafac model, including the intrinsic axis property (for details, see Harshman and Lundy, 1994, 1996).

Download

```
# load library
library(keras3)

# download Fashion MNIST data and subset training set to categories of 0, 1, 5
fmnist <- dataset_fashion_mnist()
inde <- which(fmnist$train$y %in% c(0, 1, 5))
X0 <- fmnist$train$x[inde, ,]
labels <- fmnist$train$y[inde]

# prepare labels by converting labels to class factor with levels of 0, 1, and 2
y0 <- as.factor(as.numeric(as.factor(labels)) - 1)

# subset to first 1000 images
ind <- 1:1e3
nimage <- length(ind)
Xp <- X0[ind, , ]
y <- y0[ind]

# permute the array to make the third mode the classification mode
X <- aperm(Xp, c(2, 3, 1))

# set seed
set.seed(500)

# remove rows from each image and restructure array into a list of images
nremove <- sample(3:6, nimage, replace = TRUE)
Xrag <- vector(mode = "list", length = nimage)
for (i in 1:nimage) {
  ranremove <- c(sample(1:28, nremove[i]))
```

```

image0 <- X[-ranremove, , i]
Xrag[[i]] <- image0
}




```

We plot an example of a top, trouser, and sandal.

```

# plot example of top, trouser, and sandal
par(mfrow = c(1, 3))
pdigit <- function(imat) {
  m <- t(apply(imat, 2, rev))
  image(m, col = gray(seq(0, 1, 0.05)), xaxt = "n", yaxt = "n")
}
for (i in 0:2) {pdigit(Xrag[[which(y == i)[1]]])}

```



Analysis

We load R package **cpfa** and initialize the tensor model. First, the data array is a ragged, irregular three-way array; so we set `model <- "parafac2"` to use a Parafac2 model. Second, we initialize the number of components to fit by setting `nfac <- c(2, 3)`. Third, we set `nstart <- 10` for 10 random starts. Fourth, we specify the constraint for each mode using `const`. For Parafac2, we set the third mode to have non-negative weights.

Next, we initialize classification methods. First, we use `method = c("SVM", "GBM")` to use support vector machine (SVM) and gradient boosting machine (GBM). Second, we specify that the problem is a multiclass classification problem by setting `family <- "multinomial"`. Third, we use 10-fold CV in our inner training, setting `nfolds <- 10`; and fourth, we set `nrep <- 5` to perform five outer train-test splits of our data. Fifth, we set `ratio <- 0.9` to specify that each outer training set contains a proportion of 0.9 of the full data while the outer testing set contains a proportion of 0.1. Finally, we specify ranges for tuning parameters gamma (SVM), cost (SVM), eta (GBM), maximum depth (GBM), subsample (GBM), and number of rounds (GBM), wrapping them into `parameters`.

```

# load library
library(cpfa)

# set seed

```

```

set.seed(500)

# initialize model
model <- "parafac2"
nfac <- c(2, 3)
nstart <- 10
const <- c("uncons", "uncons", "nonneg")


# initialize classification
method <- c("SVM", "GBM")
family <- "multinomial"
nfolds <- 10
nrep <- 5
ratio <- 0.9

# initialize tuning parameters
gamma <- c(0, 0.1, 1, 10); cost <- c(0.1, 1, 10, 100)
eta <- c(0.3, 0.7); max.depth <- c(1, 2)
subsample <- c(0.75, 0.9); nrounds <- c(100, 400)
parameters <- list(gamma = gamma, cost = cost, eta = eta, max.depth = max.depth,
  subsample = subsample, nrounds = nrounds)

# implement train-test splits with inner k-fold CV to optimize classification
outputR2 <- cpfa(x = Xrag, y = y, model = model, nfac = nfac, nstart = nstart,
  const = const, method = method, family = family,
  nfolds = nfolds, nrep = nrep, ratio = ratio,
  parameters = parameters, type.out = "descriptives",
  seeds = NULL, plot.out = TRUE, parallel = FALSE,
  verbose = FALSE)

```

Performance Measure



Results

We examine classification performance metrics of error (**err**) and overall accuracy (**acc**) for each model and for each classifier, averaged across train-test splits. We also examine, averaged across train-test splits, the optimal tuning parameters chosen for each classifier.

```
# examine classification performance measures - mean across train-test splits
outputR2$descriptive$mean[, 1:2]
```

```
##           err    acc
## fac.2svm 0.080 0.920
## fac.2gbm 0.100 0.900
## fac.3svm 0.080 0.920
## fac.3gbm 0.072 0.928
```

```
# examine optimal tuning parameters averaged across train-test splits
outputR2$mean.opt.tune
```


```
##   nfac alpha lambda gamma cost ntree nodesize size decay rda.alpha delta eta
## 1    2   NA     NA  6.40 26.2   NA      NA   NA   NA      NA   NA 0.54
## 2    3   NA     NA  0.64 40.6   NA      NA   NA   NA      NA   NA 0.46
##   max.depth subsample nrounds
## 1         1.6        0.84    100
## 2         1.6        0.81    100
```

We see that accuracy values are between 0.9 and 0.928. The three-component Parafac2 model with the GBM classifier performed best. In addition, we can see average optimal tuning parameters for each classifier.

Next, we use function **plotcpfa**. Because the number of levels differs among images for mode A, we cannot inspect mode A weights easily. Function **plotcpfa** automatically excludes mode A weights from plots based on the Parafac2 model. Instead, we inspect only the component weights for mode B to search for relationships between the levels of B and the Parafac2 model components.

```
# set seed
set.seed(500)

# plot heatmaps of component weights for optimal model
results2 <- plotcpfa(outputR2, nstart = 10, ctol = 1e-6, verbose = FALSE)
```



For the B mode (i.e., vertical axis), for the first component, weights are stronger between levels 11 to 14 and between 16 and 19. This first component might be capturing non-zero changes displayed at the center of the top images. The first component might help distinguish tops from trousers and sandals.

For the second component, weights are stronger between levels 6 to 23. This component could be a general component distinguishing non-zero values in the middle columns of images from zero values at the far left and far right columns. While the sandal has some non-zero values in the far left and far right columns, the top and trouser generally do not; so this second component might help distinguish the top and the trouser from the sandal.

For the third component, weights are stronger between levels 6 and 9 and between 22 and 25. This third component might be capturing zero values to the left and right of images, contrasting them with non-zero values in the middle. Interestingly, we see two negative bands between levels 11 and 14 and between 16 and 19; however, level 15 is positive. Level 15 might correspond to the empty space between the trouser legs while the negative bands might represent non-zero values in the legs. If so, this component might help distinguish trousers from tops or sandals.

These interpretations are limited. For all components, more work would need to be done to explore the estimated A mode weights (or C mode weights) to understand the components from other perspectives. Such work could clarify the components further and change the current interpretations.

C.) VesselMNIST3D from MedMNIST

The VesselMNIST3D dataset consists of 1,908 three-dimensional models of brain blood vessels based on reconstructed magnetic resonance images. Vessels exist in two categories: healthy vessels (indexed by 0) or aneurysm vessels (indexed by 1). We download the dataset manually from <https://medmnist.com/>, which gives us a .npz file. We use the Python script entitled `convert_npz_to_h5.py` to convert VesselMNIST3D from a file format of .npz to .h5. This script can be found on this project's GitHub repository. Next, we use R package **rhdf5** (Fischer, Smith, and Pau, 2025) to read the data (in .h5 file format) into R. We subset to only 300 objects from the training set, including 150 healthy vessels and 150 aneurysm vessels.

Download

```
# load library
library(rhdf5)

# import VesselMNIST3D data
h5file <- "vesselmnist3d.h5"
imag4d <- h5read(h5file, name = "train_images")
labels <- h5read(h5file, name = "train_labels")

# subset to first 300 objects
ind <- 1:300
Xp <- imag4d[ , , , ind]
y <- as.factor(as.numeric(labels))[ind]

# convert array to double
storage.mode(Xp) <- "double"

# remove full dataset
rm(imag4d)
```


We use R package **plotly** (Sievert, 2020) to generate a point cloud representation of a healthy blood vessel.

```
# load library
library(plotly)

# prepare data from vessel 205 for plotting
voxel_df <- as.data.frame(which(Xp[, , , 205] > 0, arr.ind = TRUE))
colnames(voxel_df) <- c("x", "y", "z")

# generate 3D plot
scene1 <- list(title = "", showticklabels = FALSE, showgrid = TRUE,
               showline = FALSE, zeroline = FALSE)
plot_ly(voxel_df, x = ~x, y = ~y, z = ~z, type = "scatter3d", mode = "markers",
        marker = list(size = 3.5, color = "blue", opacity = 0.3)) %>%
  layout(title = "3D Blood Vessel Example - Point Cloud Representation",
        scene = list(xaxis = scene1, yaxis = scene1, zaxis = scene1))
```

3D Blood Vessel Example - Point Cloud Representation



Analysis

We load R package **cpfa** and initialize the tensor model. First, the data array is a regular four-way array; so we set `model <- "parafac"` to use a Parafac model. Second, we initialize the number of components to fit by setting `nfac <- c(2, 3)`. Third, we set `nstart <- 10` for 10 random starts. Fourth, we specify the constraint for each mode using `const`. For Parafac, we set the fourth mode to have non-negative weights.

Next, we initialize classification methods. First, we use `method = c("NN", "RDA")` to use feed-forward neural networks (NN) and regularized discriminant analysis (RDA). Second, we specify that the problem is a binary classification problem by setting `family <- "binomial"`. Third, we use 10-fold CV in our inner training, setting `nfolds <- 10`; and fourth, we set `nrep <- 5` to perform five outer train-test splits of our data. Fifth, we set `ratio <- 0.9` to specify that each outer training set contains a proportion of 0.9 of the full data while the outer testing set contains a proportion of 0.1. Finally, we specify ranges for tuning parameters alpha (RDA), delta (RDA), size (NN), and decay (NN), wrapping them into `parameters`.

```
# load library
library(cpfa)

# set seed
set.seed(500)

# initialize model
model <- "parafac"
nfac <- c(2, 3)
nstart <- 10
const <- c("uncons", "uncons", "uncons", "nonneg")


# initialize classification
method <- c("NN", "RDA")
family <- "binomial"
nfolds <- 10
nrep <- 5
ratio <- 0.9

# initialize tuning parameters
rda.alpha <- seq(0, 0.999, length = 4); delta <- c(0, 0.1, 1, 2)
```

```
size <- c(1, 2, 4, 8); decay <- c(0.001, 0.01, 0.1, 1)
parameters <- list(rda.alpha = rda.alpha, delta = delta, size = size,
                  decay = decay)

# implement train-test splits with inner k-fold CV to optimize classification
outputR3 <- cpfa(x = Xp, y = y, model = model, nfac = nfac, nstart = nstart,
               const = const, method = method, family = family,
               nfolds = nfolds, nrep = nrep, ratio = ratio,
               parameters = parameters, type.out = "descriptives",
               seeds = NULL, plot.out = TRUE, parallel = FALSE,
               verbose = FALSE)
```

Performance Measure



Results

We examine classification performance metrics of error (`err`) and overall accuracy (`acc`) for each model and for each classifier, averaged across train-test splits. We also examine, averaged across train-test splits, the optimal tuning parameters chosen for each classifier.

```
# examine classification performance measures - mean across train-test splits
outputR3$descriptive$mean[, 1:2]
```

```
##           err           acc
## fac.2nn  0.3466667 0.6533333
## fac.2rda 0.3400000 0.6600000
## fac.3nn  0.3133333 0.6866667
## fac.3rda 0.2933333 0.7066667
```

```
# examine optimal tuning parameters averaged across train-test splits
outputR3$mean.opt.tune
```

```
##   nfac alpha lambda gamma cost ntree nodesize size  decay rda.alpha delta eta
```


```
## 1      2      NA      NA      NA      NA      NA      NA      3.6 0.0082      0.5328 0.02 NA
## 2      3      NA      NA      NA      NA      NA      NA      NA      5.6 0.0226      0.5328 0.04 NA
##      max.depth subsample nrounds
## 1      NA      NA      NA
## 2      NA      NA      NA
```

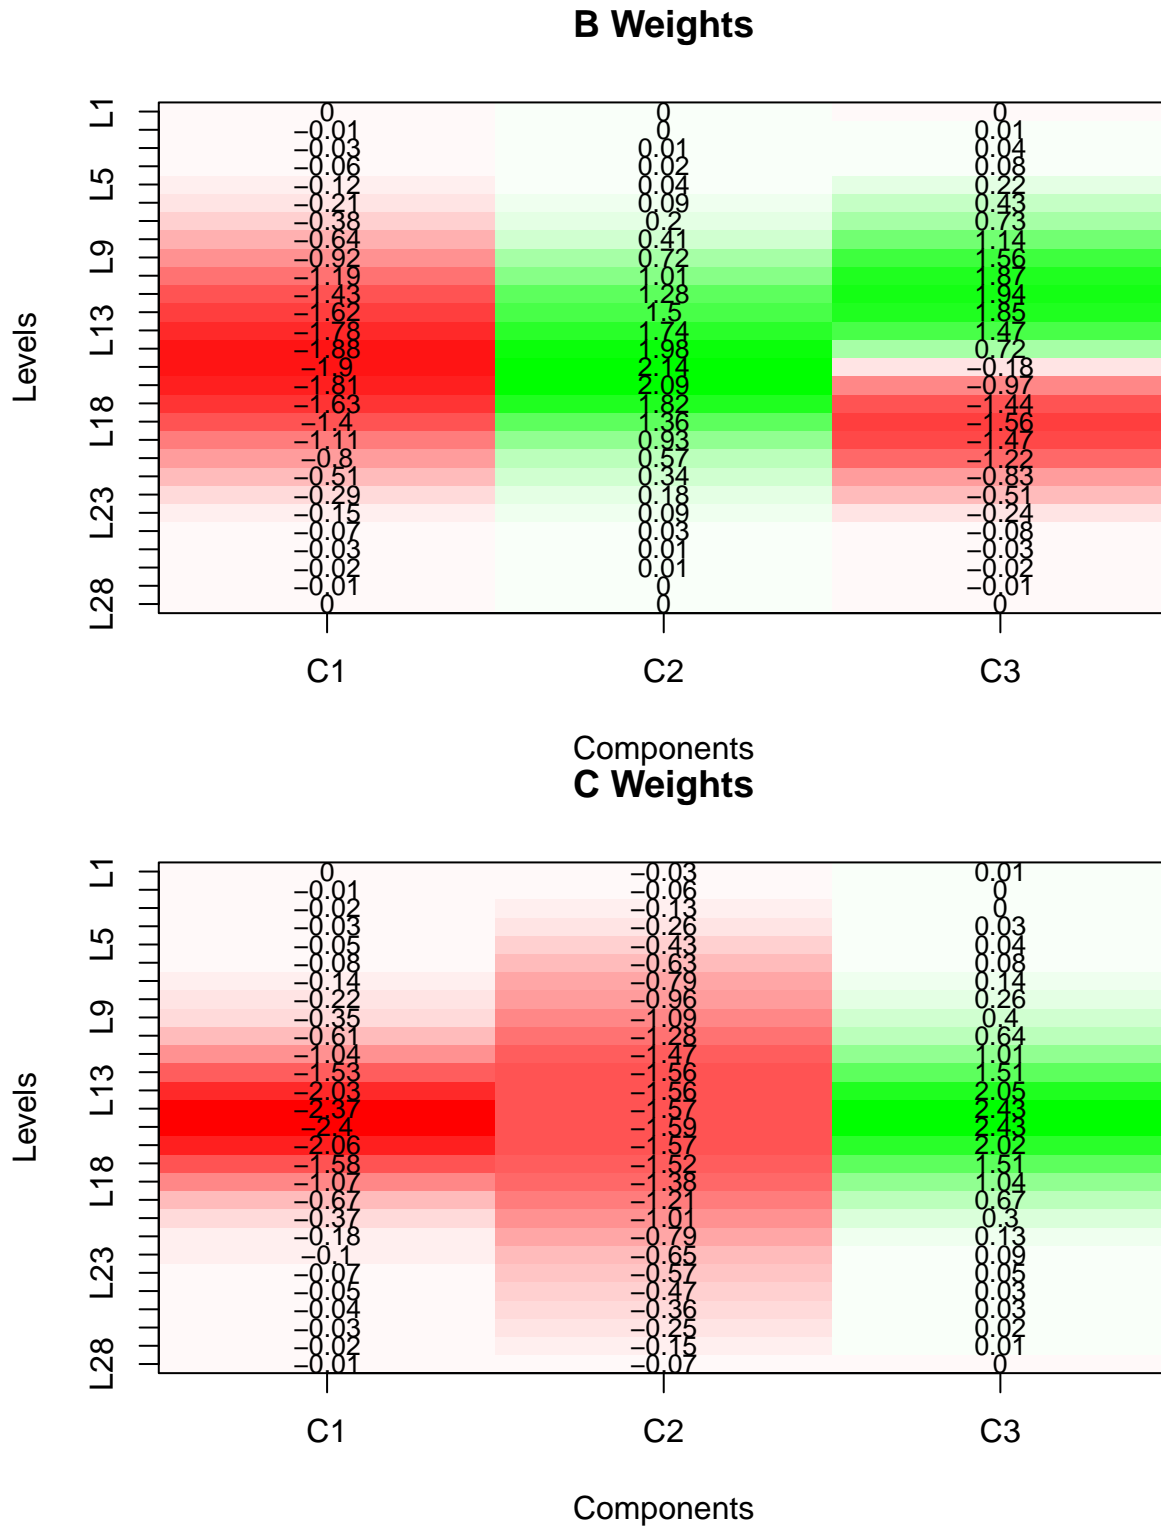
We see that accuracy values are between 0.65 and 0.71. The three-component Parafac model with the RDA classifier performed best. In addition, we can see average optimal tuning parameters for each classifier.

Next, we use function `plotcpfa`. Looking at the plots, we can search for relationships between the levels of mode A, B, or C and the model components. In this case, the three modes correspond to the three dimensions of the vessels.

```
# set seed
set.seed(500)

# plot heatmaps of component weights for optimal model
results3 <- plotcpfa(outputR3, nstart = 10, ctol = 1e-6, verbose = FALSE)
```





For now, we omit interpretations of the three components, given the difficulty of visualizing each blood vessel object. For more information about this dataset and classification work using it, interested parties might consider exploring Yang, Xia, Kin, and Igarashi (2020); Yang et al. (2023); and Yang, Shi, and Ni (2021).

References

- Asisgress, M. (2025). `cpfa`: Classification with Parallel Factor Analysis. R package version 1.2-2, <https://CRAN.R-project.org/package=cpfa>.
- Fischer, B., Smith, M., and Pau, G. (2025). `rhdf5`: R Interface to HDF5. R package version 2.52.1, <https://bioconductor.org/packages/rhdf5>.
- Harshman, R. (1970). Foundations of the PARAFAC procedure: Models and conditions for an “explanatory” multi-modal factor analysis. *UCLA working papers in phonetics*, 16(1), 84.
- Harshman, R. (1972). PARAFAC2: Mathematical and technical notes. *UCLA working papers in phonetics*, 22(3044), 122215.
- Harshman, R. and Lundy, M. (1994). PARAFAC: Parallel factor analysis. *Computational Statistics and Data Analysis*, 18(1), 39-72.
- Harshman, R. and Lundy, M. (1996). Uniqueness proof for a family of models sharing features of Tucker’s three-mode factor analysis and PARAFAC/CANDECOMP. *Psychometrika*, 61(1), 133-154.
- Helwig, N. (2025). `CMLS`: Constrained Multivariate Least Squares. R package version 1.1, <https://CRAN.R-project.org/package=CMLS>.
- Helwig, N. (2025). `multiway`: Component Models for Multi-Way Data. R package version 1.0-7, <https://CRAN.R-project.org/package=multiway>.
- Irizarry, R., Gill, A. (2025). `dslabs`: Data Science Labs. R package version 0.9.0, <https://CRAN.R-project.org/package=dslabs>.
- Kalinowski, T., Allaire, J., and Chollet, F. (2025). `keras3`: R Interface to ‘Keras’. R package version 1.4.0, <https://CRAN.R-project.org/package=keras3>.
- LeCun, Y., Bottou, L., Bengio, Y., and Haffner, P. (2002). Gradient-based learning applied to document recognition. *Proceedings of the IEEE*, 86(11), 2278–2324.
- LeCun, Y., Cortes, C., and Burges, C. (1998). The MNIST database of handwritten digits. <http://yann.lecun.com/exdb/mnist/>.
- R Core Team (2025). R: A Language and Environment for Statistical Computing. R Foundation for Statistical Computing, Vienna, Austria. <https://www.R-project.org/>.
- Sievert, C. (2020). Interactive web-based data visualization with R, plotly, and shiny. Chapman and Hall/CRC.
- Xiao, H., Rasul, K., and Vollgraf, R. (2017). Fashion-mnist: a novel image dataset for benchmarking machine learning algorithms. *arXiv preprint arXiv:1708.07747*.
- Yang, J., Shi, R., and Ni, B. (April 2021). Medmnist classification decathlon: A lightweight automl benchmark for medical image analysis. In *2021 IEEE 18th International Symposium on Biomedical Imaging (ISBI)* (pp. 191-195). IEEE.
- Yang, J., Shi, R., Wei, D., Liu, Z., Zhao, L., Ke, B., . . . and Ni, B. (2023). Medmnist v2-a large-scale lightweight benchmark for 2d and 3d biomedical image classification. *Scientific Data*, 10(1), 41.
- Yang, X., Xia, D., Kin, T., and Igarashi, T. (2020). Intra: 3d intracranial aneurysm dataset for deep learning. In *Proceedings of the IEEE/CVF conference on computer vision and pattern recognition* (pp. 2656-2666).